

# Reinforcement Learning for Addition via Recursion

Mathematical Reasoning Project

May 10, 2025

## Abstract

This report presents an implementation of a Reinforcement Learning (RL) agent designed to learn and execute the recursive definition of addition. The agent learns to transform expressions like  $a+b$  into  $(a+(b-1))+1$  recursively until reaching the base case  $a+0=a$ . We demonstrate the agent's capability by solving the addition problem  $5+8$  and visualizing the solution path as a knowledge graph in Neo4j. This implementation showcases how RL can be applied to mathematical reasoning tasks.

## 1 Introduction

Addition is typically taught as a basic operation, but it can be formally defined using recursion:

$$a + b = \begin{cases} a & \text{if } b = 0 \quad (\text{Base Case}) \\ (a + (b - 1)) + 1 & \text{if } b > 0 \quad (\text{Recursive Case}) \end{cases} \quad (1)$$

This recursive definition provides a systematic way to reduce complex addition problems to simpler ones. Our implementation uses Reinforcement Learning to teach an agent to apply this recursive definition correctly.

## 2 Reinforcement Learning Components

### 2.1 Key RL Concepts

Reinforcement Learning is a machine learning paradigm where an agent learns to make decisions by interacting with an environment. The key components are:

- **Agent:** The decision-maker that learns from experience
- **Environment:** The world with which the agent interacts
- **State:** The current situation in the environment

- **Action:** A decision made by the agent
- **Reward:** Feedback signal indicating the quality of an action
- **Policy:** The agent’s strategy for selecting actions
- **Value Function:** Estimation of future rewards from a state
- **Q-Function:** Estimation of future rewards from a state-action pair

## 2.2 Q-Learning Algorithm

Our implementation uses Q-learning, a model-free RL algorithm that learns the value of an action in a particular state. The Q-function is updated using:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \cdot [r + \gamma \cdot \max_{a'} Q(s', a') - Q(s, a)] \quad (2)$$

where:

- $Q(s, a)$  is the value of taking action  $a$  in state  $s$
- $\alpha$  is the learning rate
- $r$  is the reward received
- $\gamma$  is the discount factor for future rewards
- $s'$  is the next state
- $\max_{a'} Q(s', a')$  is the maximum Q-value for the next state

## 3 Implementation Components

### 3.1 Environment

The environment (`AdditionRecursionEnv`) represents the addition problem and tracks:

- Current mathematical expression state
- Target result
- Available actions
- Step count and termination conditions

### 3.2 State Representation

States are represented by the `MathExpressionState` class, which contains:

- The current mathematical expression (e.g., `"5 + 8"`, `"(5 + 7) + 1"`)
- Step count for tracking progress

### 3.3 Actions

The agent can perform the following actions:

- **decompose**: Transform " $a + b$ " to " $(a + (b-1)) + 1$ "
- **further\_decompose**: Further decompose a nested expression
- **apply\_base\_case**: Apply the rule  $(a + 0) = a$
- **increment**: Perform an increment operation (e.g.,  $(3 + 1)$  to 4)
- **calculate**: Calculate the final result
- **finish**: Conclude the calculation

### 3.4 Reward Structure

The reward structure guides the agent toward correct solutions:

- +1.0 for correct decomposition steps
- +1.0 for correctly applying the base case
- +1.0 for correct increment operations
- +5.0 for reaching the correct final result
- -1.0 for invalid actions
- -5.0 for finishing with an incorrect result

### 3.5 Agent

The `QLearningAgent` implements:

- Q-table for storing state-action values
- Epsilon-greedy policy for action selection
- Q-learning update rule
- Exploration rate decay
- Action history tracking

### 3.6 Knowledge Graph Builder

The `KnowledgeGraphBuilder` constructs a graph representation of the solution:

- Nodes represent mathematical expressions (states)
- Edges represent transformations (actions) between states
- The graph is exported as triplets for Neo4j storage

## 4 Training Process

The training process involves:

1. Initializing the environment with the problem "5 + 8"
2. Training the agent for a specified number of episodes
3. For each episode:
  - Reset the environment
  - Agent selects actions using epsilon-greedy policy
  - Environment provides rewards and next states
  - Agent updates Q-values using the Q-learning update rule
4. Tracking the best solution found during training
5. Building a knowledge graph from the best solution
6. Storing the graph in Neo4j

## 5 Example: Solving 5 + 8

For the addition problem 5 + 8, the ideal solution path is:

$$5 + 8 = (5 + 7) + 1 \quad (3)$$

$$= ((5 + 6) + 1) + 1 \quad (4)$$

$$= (((5 + 5) + 1) + 1) + 1 \quad (5)$$

$$= ((((5 + 4) + 1) + 1) + 1) + 1 \quad (6)$$

$$= ((((((5 + 3) + 1) + 1) + 1) + 1) + 1) + 1 \quad (7)$$

$$= (((((((5 + 2) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (8)$$

$$= (((((((((5 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (9)$$

$$= ((((((((((5 + 0) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (10)$$

$$= ((((((((((5 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (11)$$

$$= ((((((((((6 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (12)$$

$$= ((((((((((7 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (13)$$

$$= ((((((((((8 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (14)$$

$$= ((((((((((9 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (15)$$

$$= ((((((((((10 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (16)$$

$$= ((((((((((11 + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1) + 1 \quad (17)$$

$$= 12 + 1 \quad (18)$$

$$= 13 \quad (19)$$

## 6 Knowledge Graph Representation

The solution is represented as a knowledge graph in Neo4j:

- **Nodes:** Mathematical expressions at each step
- **Edges:** Actions taken to transform expressions
- **Properties:** Step information, action types, and rewards

This graph provides a visual representation of the reasoning process, showing how the agent decomposes the problem and builds up to the solution.

## 7 Conclusion

This implementation demonstrates how Reinforcement Learning can be applied to mathematical reasoning tasks. The agent successfully learns to apply the recursive definition of addition and solve problems like  $5 + 8$ . The knowledge graph representation provides insight into the reasoning process and can be used for educational purposes.

Future work could include:

- Extending the approach to other mathematical operations (subtraction, multiplication)
- Implementing more complex mathematical reasoning tasks
- Improving the agent's learning efficiency
- Developing a user interface for interactive exploration of the solution paths